


ServiceNow AI Intelligence Platform

Demo Talking Points | June 2026 | Capgemini AI & Data Practice

Platform Overview

Transforms raw ServiceNow incident exports into operational intelligence across the full incident lifecycle. Built on **FastAPI + React**, trained on **1,200 incidents (Nov 2025 – Mar 2026)**. Seven live modules from real-time monitoring through to conversational AI.



Talking Points

- Single-screen situational awareness — 429 open incidents, no manual report-pulling needed.
- P1 count (21) is above the 15-ticket alert threshold; live operational risk flagged on-screen.
- SLA compliance at 36.2% vs 90% target — the heat map pinpoints exactly which group and priority is dragging it down.
- 97.8 h MTTR is the key lever; every hour cut here directly lifts SLA compliance.

Chart Rationale

Chart	What it shows	Key insight
Incident Volume by Group	Stacked bars by state per group	WEB-L2 carries highest load (387). GlobalService Desk most In-Progress — bottleneck.
By Category (Donut)	% split by category	Application Access 38.7% — provisioning/access process gap likely root cause.
SLA KPI Gauge	Overall compliance + per-priority	Instant red/amber/green signal; P1 & P2 non-compliance most business-critical.
Priority Heat Map	P1–P4 count per group	WEB-L2 has 235 P4 tickets — low-priority overload can mask P1 escalation risk.
Reopen Tracker	Monthly reopen trend + top tickets	Reopens indicate first-time resolution failure — direct link to MTTR inflation.



Talking Points

- Answers the QBR question: 'are we improving?' — month-by-month trend across all key metrics.
- SLA compliance 36.79% — area chart makes deterioration periods immediately visible for exec conversations.
- MTTR by Group lines cross over time — shows which teams improved and which need support, without a spreadsheet.
- 6-period ARIMA forecast lets teams pre-position resources before volume spikes hit capacity.
- Root-cause treemap drills category → subcategory; surfaces the single subcategory driving the most volume.

Chart Rationale

Chart	What it shows	Key insight
-------	---------------	-------------

Incident Volume Over Time	Stacked bar by group + trend line	Volume broadly stable; trend overlay (+10.8%) flags acceleration risk.
SLA Compliance Trend	Met vs breached per period (area)	Compliance 36.79% — area shading makes breached periods visually dominant.
MTTR by Group	Avg resolution hours per group	Crossed lines show PACS-L2 improving; Global Service Desk worsening — coaching
Priority Trend	P1-P4 stacked bars per period	P4 dominates — useful to show workload is not all critical incidents.
Resolution Heatmap	Tickets resolved by day x hour	Peak resolution: weekday business hours. Weekend gaps = on-call coverage opport
Reassignment Analysis	Avg reassignments + scatter	High reassignment groups have longer MTTR — scatter confirms the correlation.

M3	Smart Triage	AI auto-classification · Trained on 1,200 incidents
----	--------------	---

Talking Points

- L1 agents currently spend 3-5 min categorising each ticket manually — M3 does it in under 2 seconds.
- Category: 100% accuracy · Routing: 100% · Priority: 59% (inherently subjective; used as a hint, not a decision).
- 'Top 5 similar past incidents' surfaces proven resolution paths instantly — cuts handle time and reduces reopens.
- Optional LLM Resolution Hint grounds the AI answer in your own historical data — no hallucination risk.

Model Accuracy

Model	Technique	Accuracy	Purpose
Category classifier	TF-IDF + Linear SVC	100%	Assign correct incident type
Group router	TF-IDF + Linear SVC	100%	Assign to correct assignment group
Priority predictor	Risk classification	59%	Flag urgency for L1 review
Similarity engine	Cosine similarity	—	Surface 5 nearest historical tickets

M4/M5	Intelligent Routing & NL Chatbot	Auto-assignment ML · Conversational AI
-------	----------------------------------	--

M4 — Intelligent Routing

- Predicts the best assignment group from free-text — eliminates manual routing decisions end-to-end.
- Directly reduces reassignment count, which is the strongest leading indicator of MTTR inflation.

M5 — Conversational Chatbot

- Ask plain-English questions: 'Which group has the worst SLA this month?' — no dashboard navigation required.
- RAG architecture: answers grounded in live incident data, not generic LLM training — zero hallucination on operational facts.
- Designed for SDMs and executives who need quick answers in meeting contexts.

M6	SDM Scorecard	AM metric compliance · Agent performance KPIs
----	---------------	---

Talking Points

- Single-screen Green / Amber / Red view of all AM metric thresholds — replaces manual weekly reporting.
- Agent Breakdown tab enables individual coaching conversations backed by data, not perception.
- Monthly Trend tab demonstrates continuous service improvement trajectory to the client in seconds.
- Fully filter-driven — drill to a specific group, priority, or date range without touching a spreadsheet.

Summary

Executive Value Summary

One-line ROI per module

Module	Value Delivered	Status
M1 — Monitoring	Real-time queue health & SLA tracking, no manual reporting	Live
M2 — Trends	6-month insight & 6-period demand forecast for capacity planning	Live
M3 — Smart Triage	Auto-classify in <2s; surfaces similar resolutions to cut MTTR	Live
M4 — Routing	Eliminate manual group assignment with ML; reduce reassignments	Live
M5 — Chatbot	Natural-language Q&A over incident data, RAG-grounded	Live
M6 — Scorecard	SDM-level AM metric compliance & individual agent KPIs	Live
M7 — Data Upload	CSV/Excel ingestion with dedup, sanitisation & column mapping	Live